

# Productivity meets Performance Julia on A64FX

Mosè Giordano  
[m.giordano@ucl.ac.uk](mailto:m.giordano@ucl.ac.uk)



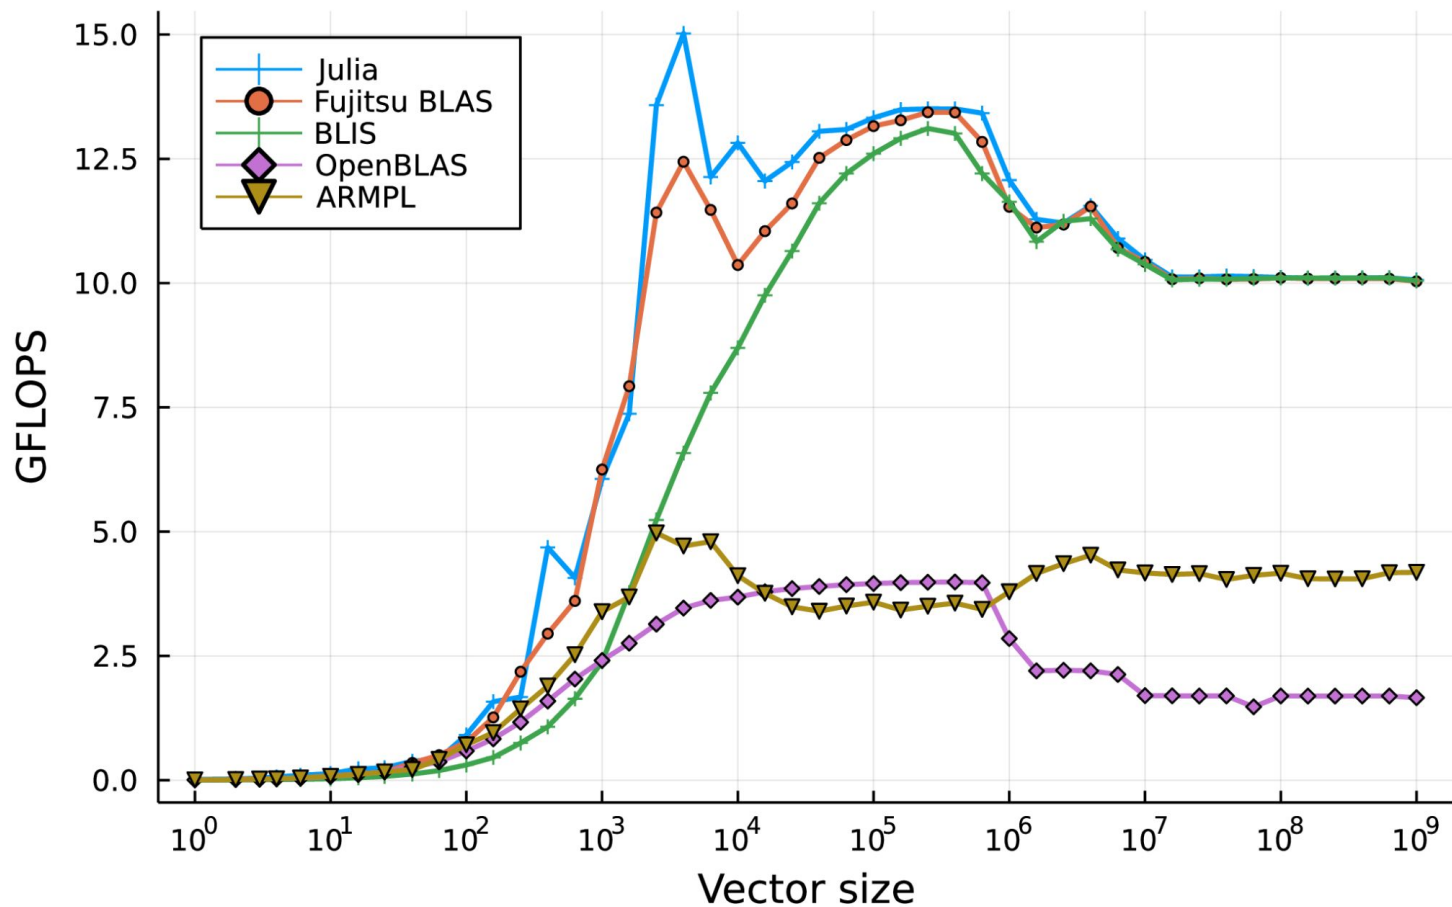
# Level 1 BLAS showdown

```
function axpy!(a, x, y)
    @simd for i in eachindex(x, y)
        @inbounds y[i] = muladd(a, x[i], y[i])
    end
    return y
end
```

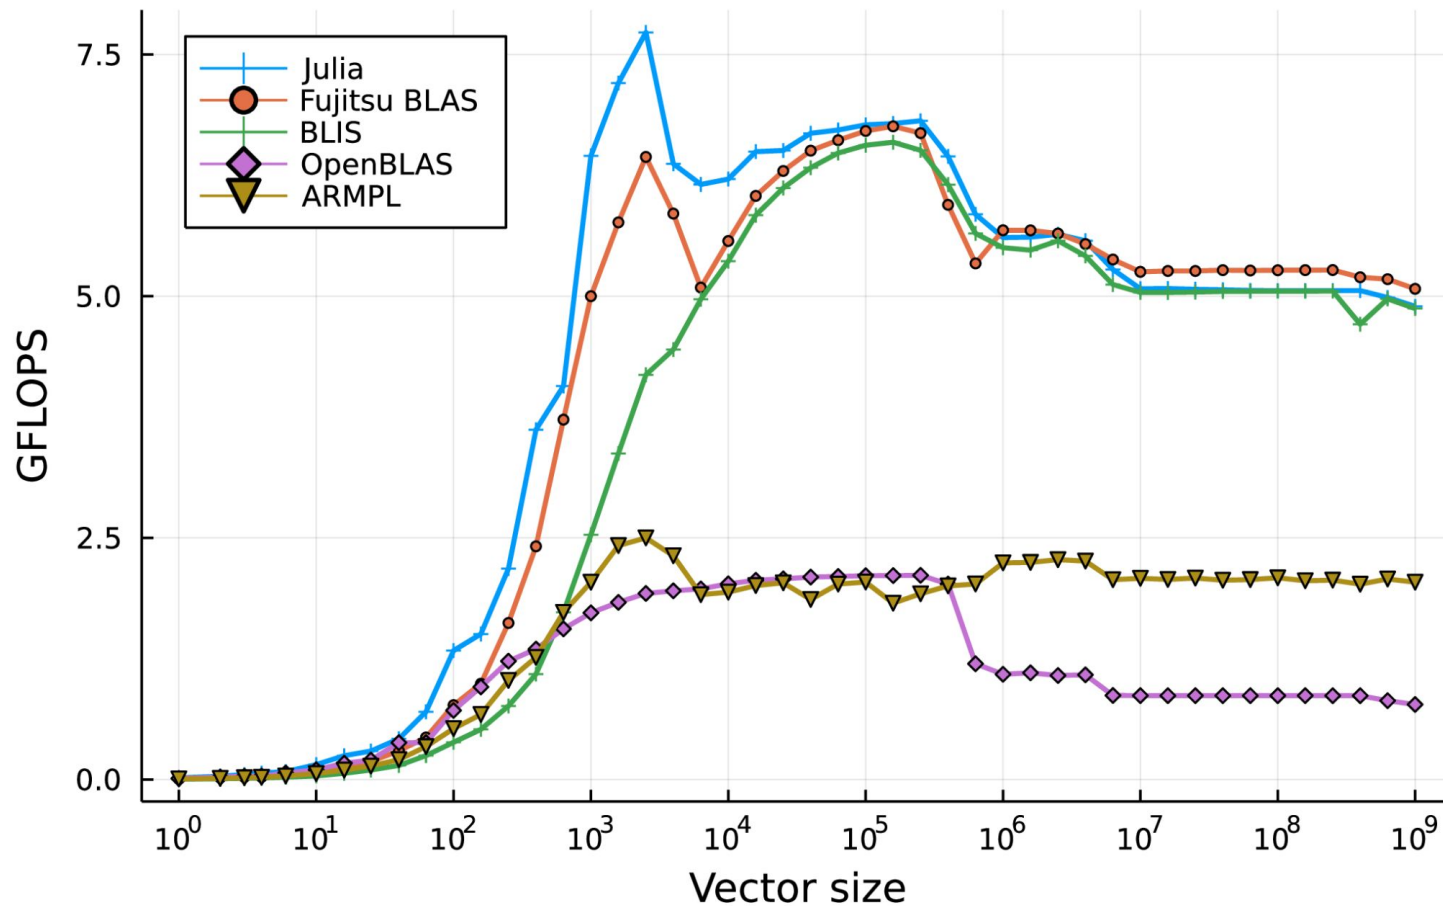
vs

```
LinearAlgebra.BLAS.axpy!(a, x, y)
```

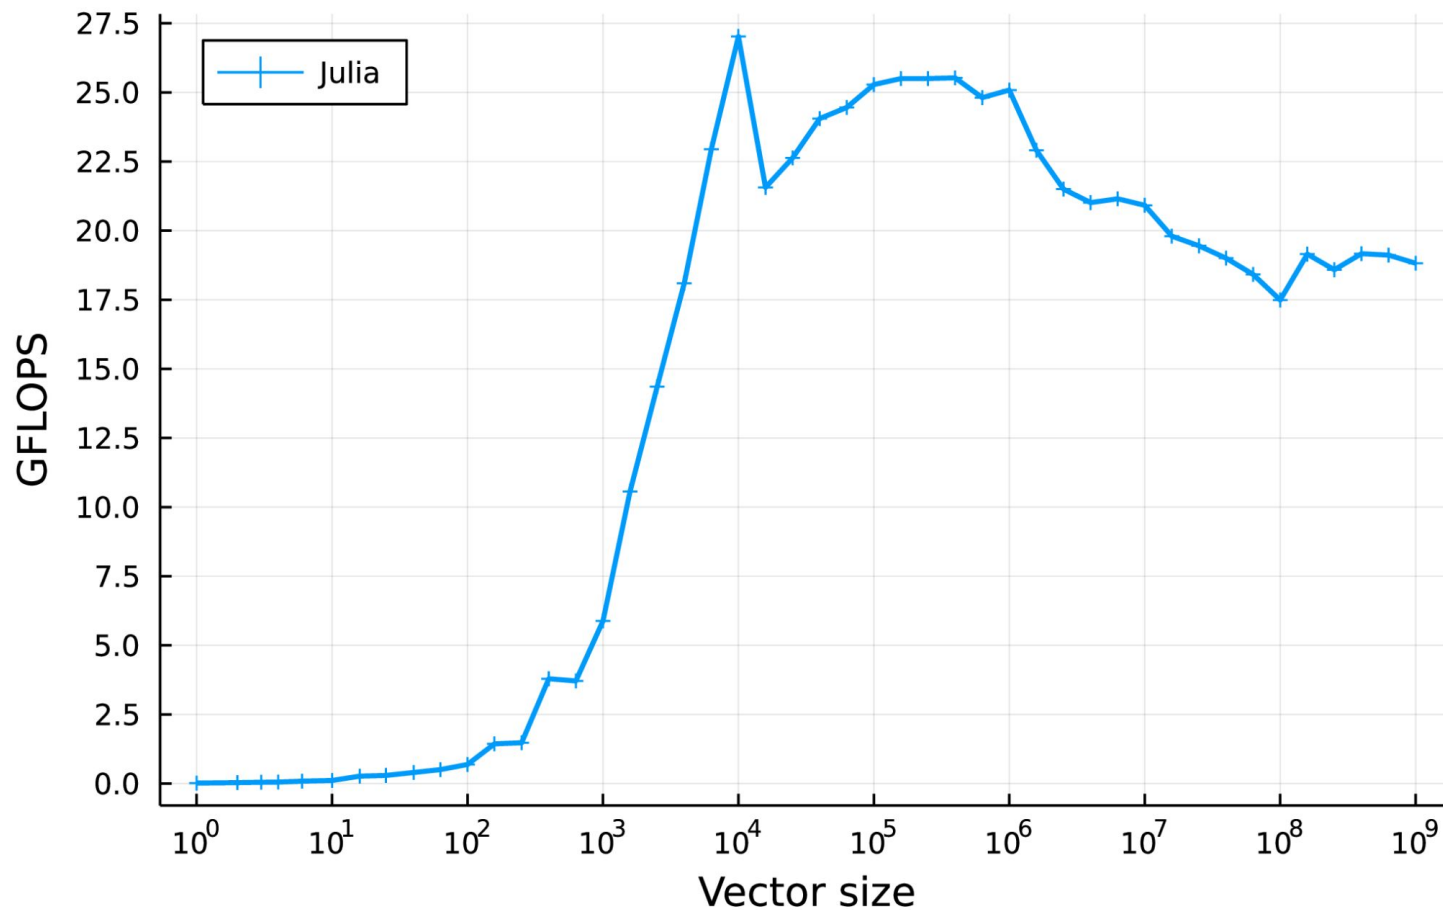
# axpy (single precision)



# axpy (double precision)

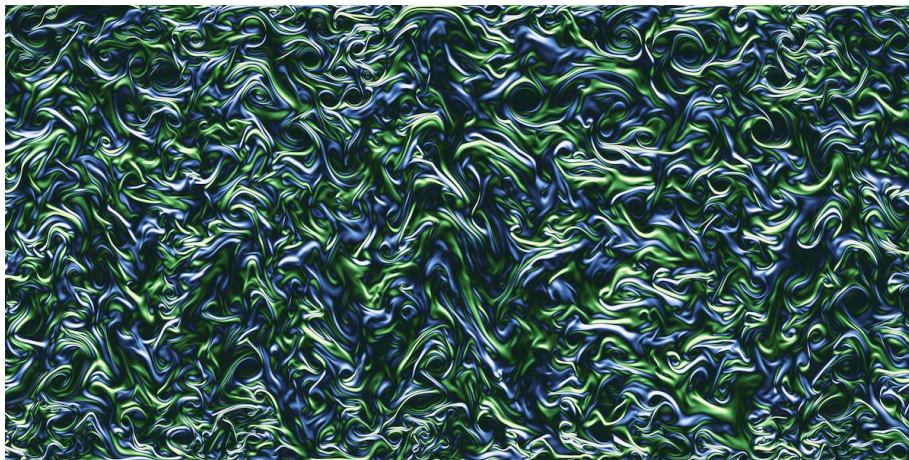


# axpy (half precision)



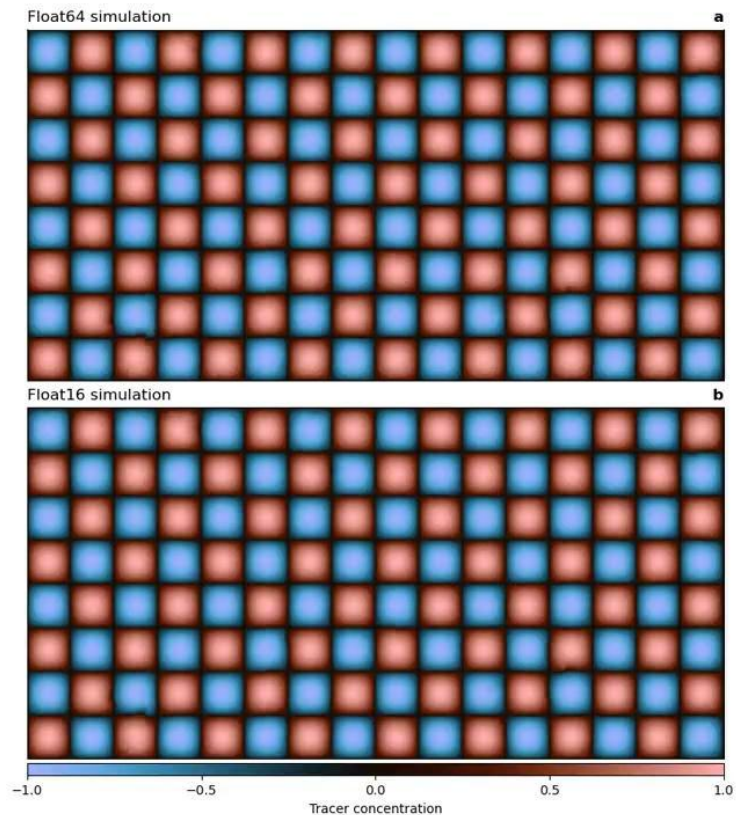
# ShallowWaters.jl

- Open-Source CFD code written in Julia
- Type-agnostic/Type-flexible
  - Compensated summation for low-precision
- ~4x speedup with Float16 and 2x speedup with Float32 over Float64
- Qualitative results equivalent between Float64 and Float16

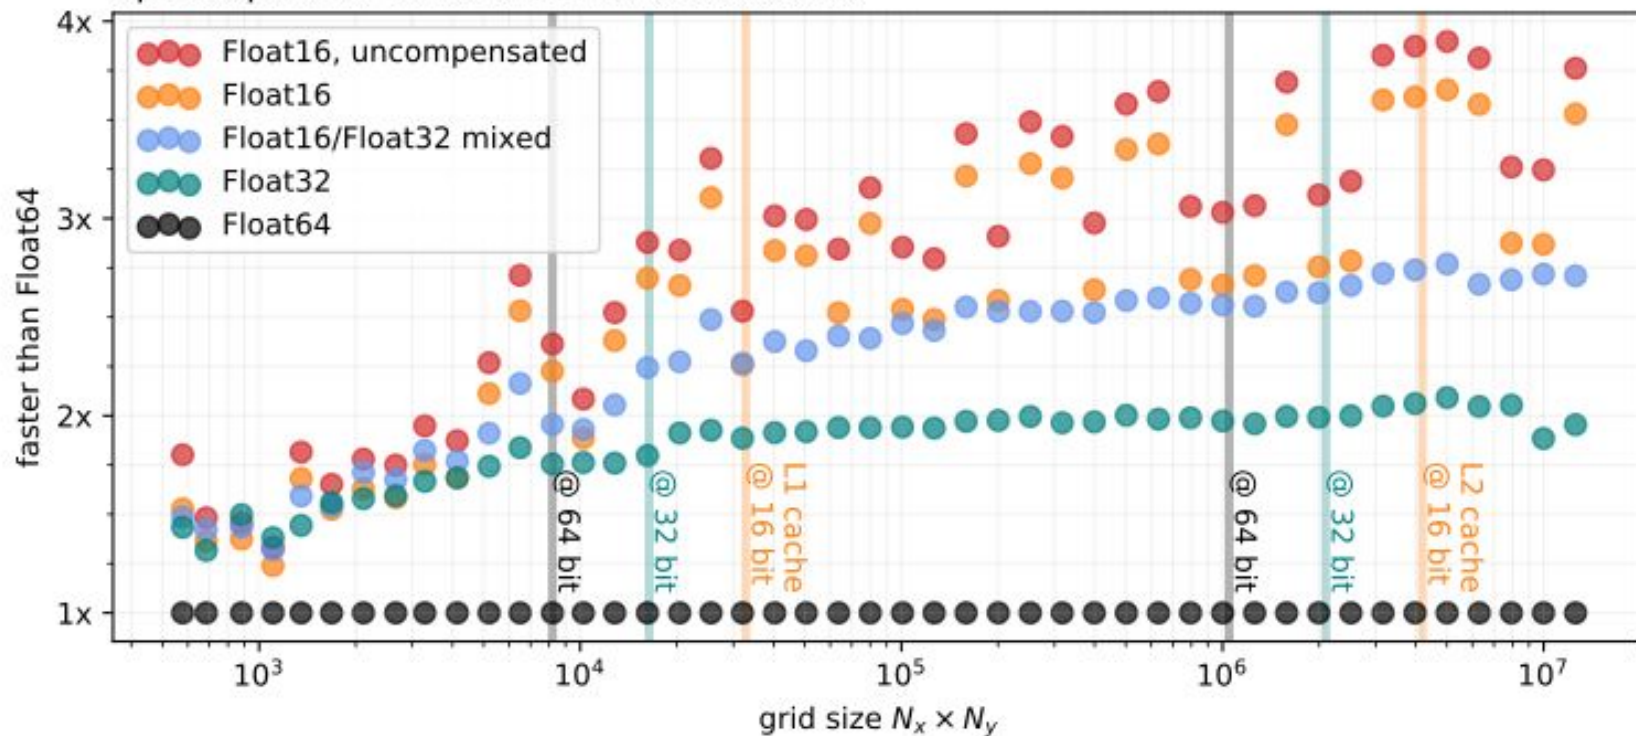




# ShallowWaters.jl — Fidelity comparison



Speedups with 16-bit arithmetic on A64FX





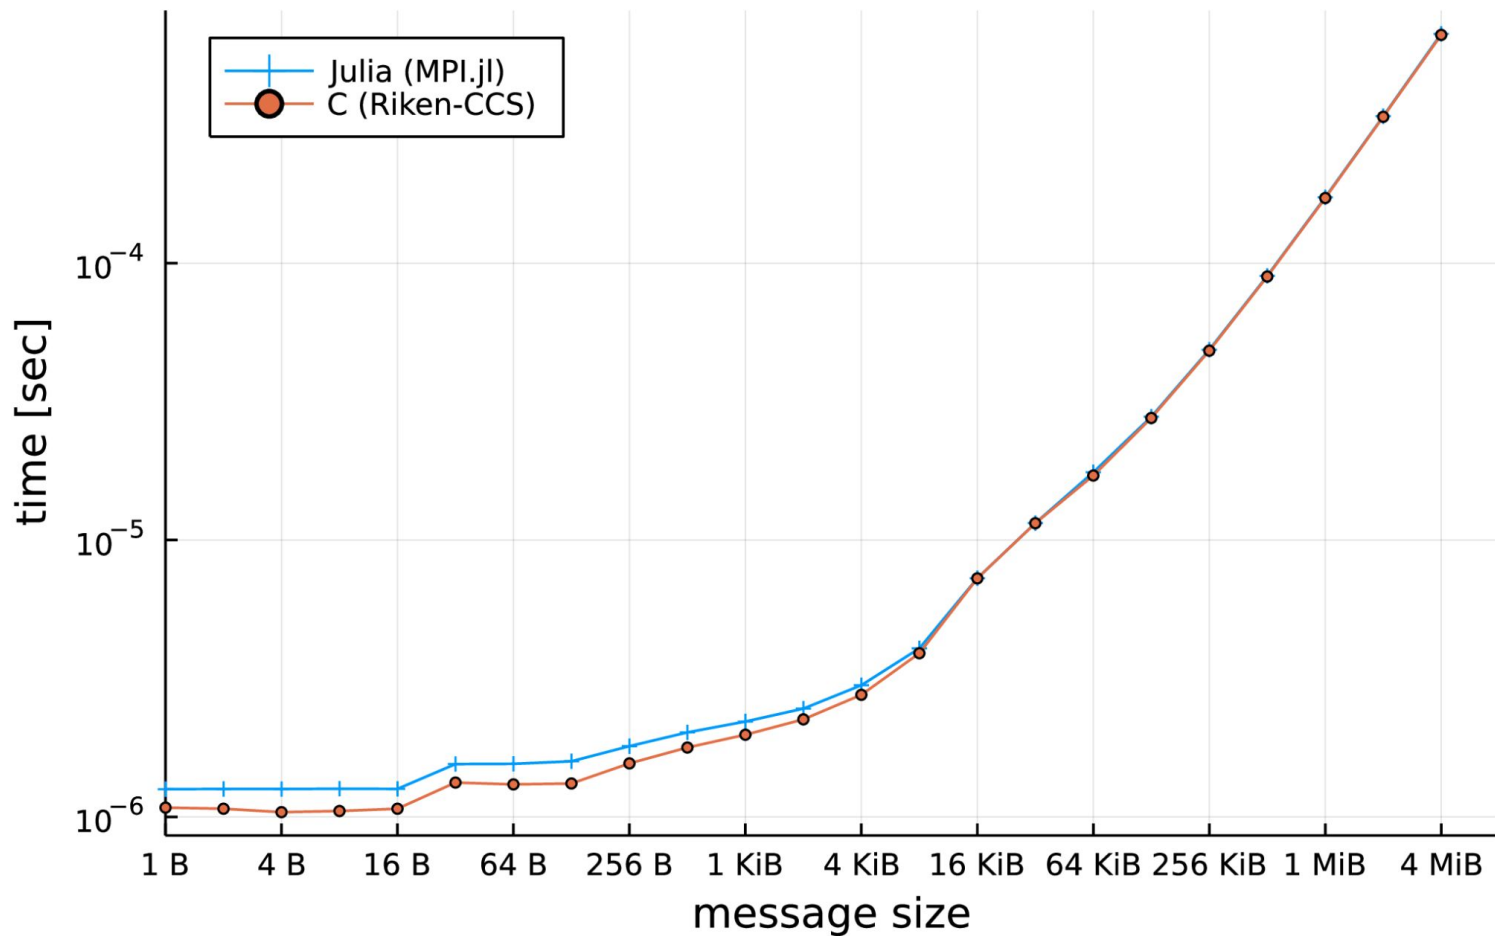
# MPI.jl

- Low-level access to MPI
- High-level convenience wrappers
- Deals with MPI ABI

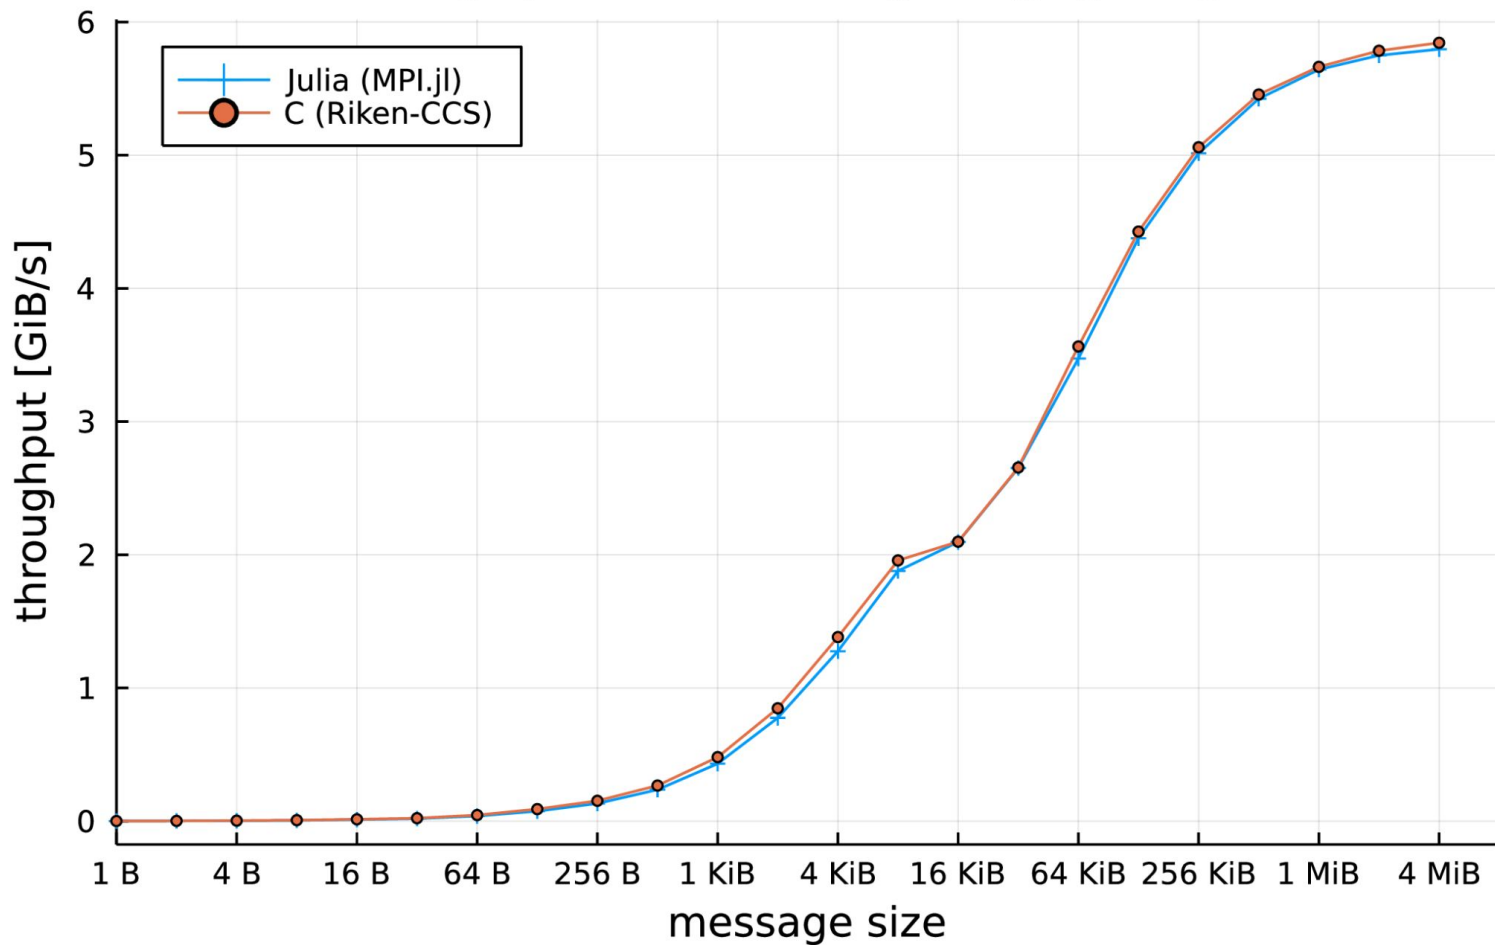
One of the oldest Julia packages (2012)

```
function pingpong(T::Type, bufsize::Int,
                  iters::Int, comm::MPI.Comm)
    rank = MPI.Comm_rank(comm)
    buffer = zeros(T, bufsize)
    tag = 0
    MPI.Barrier(comm)
    tic = MPI.Wtime()
    for i in 1:iters
        if iszero(rank)
            MPI.Send(buffer, comm; dest=1, tag)
            MPI.Recv!(buffer, comm; source=1, tag)
        elseif isone(rank)
            MPI.Recv!(buffer, comm; source=0, tag)
            MPI.Send(buffer, comm; dest=0, tag)
        end
    end
    toc = MPI.Wtime()
    return (toc - tic) / iters
end
```

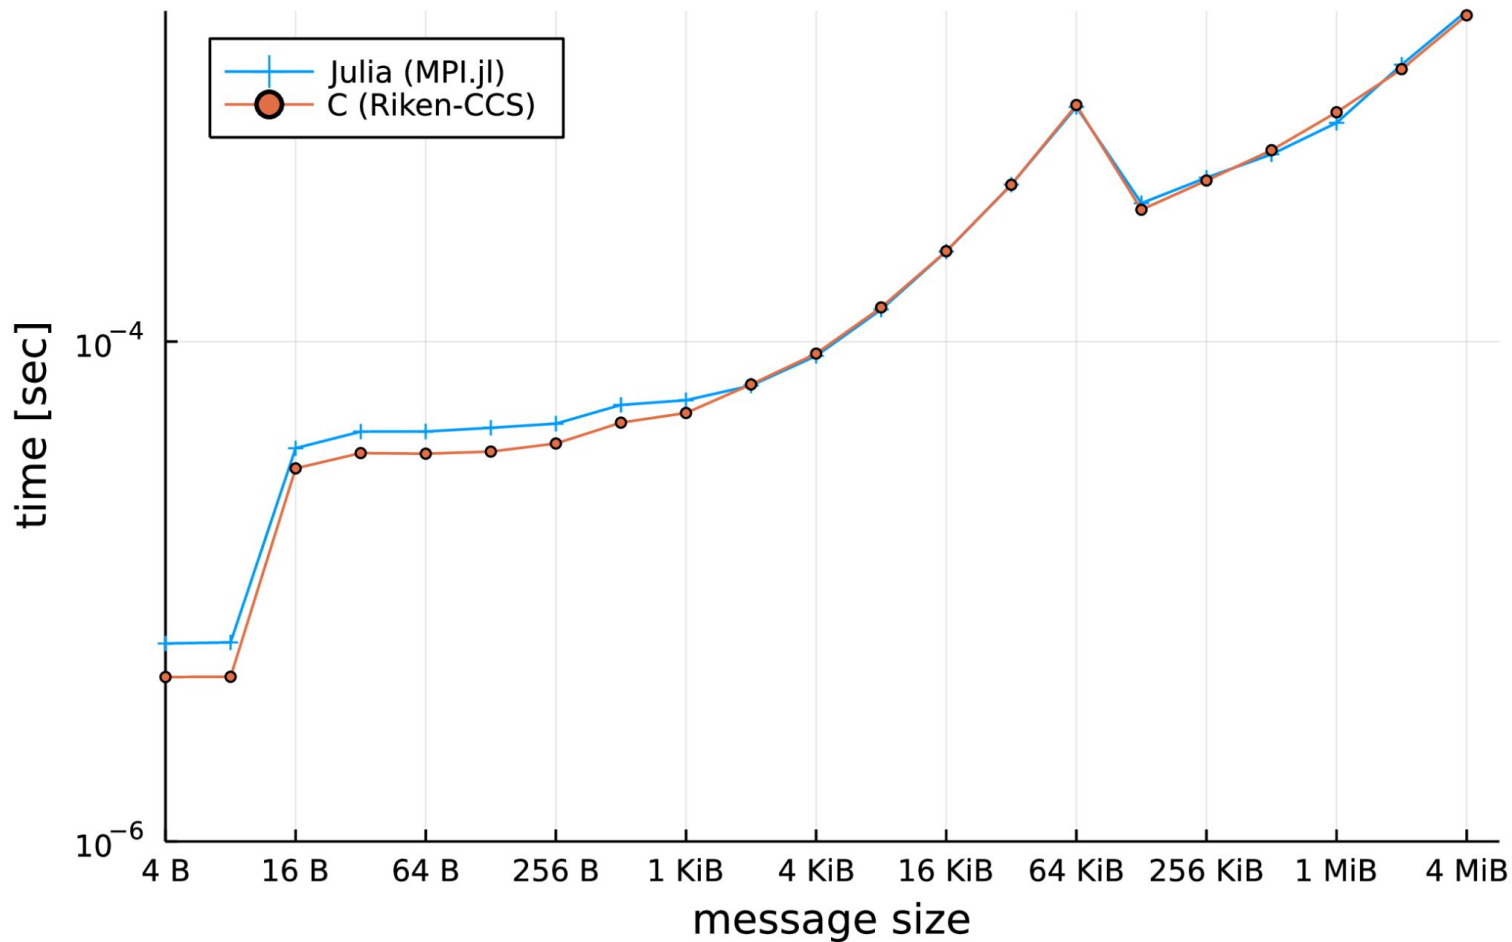
# Latency of MPI PingPong @ Fugaku



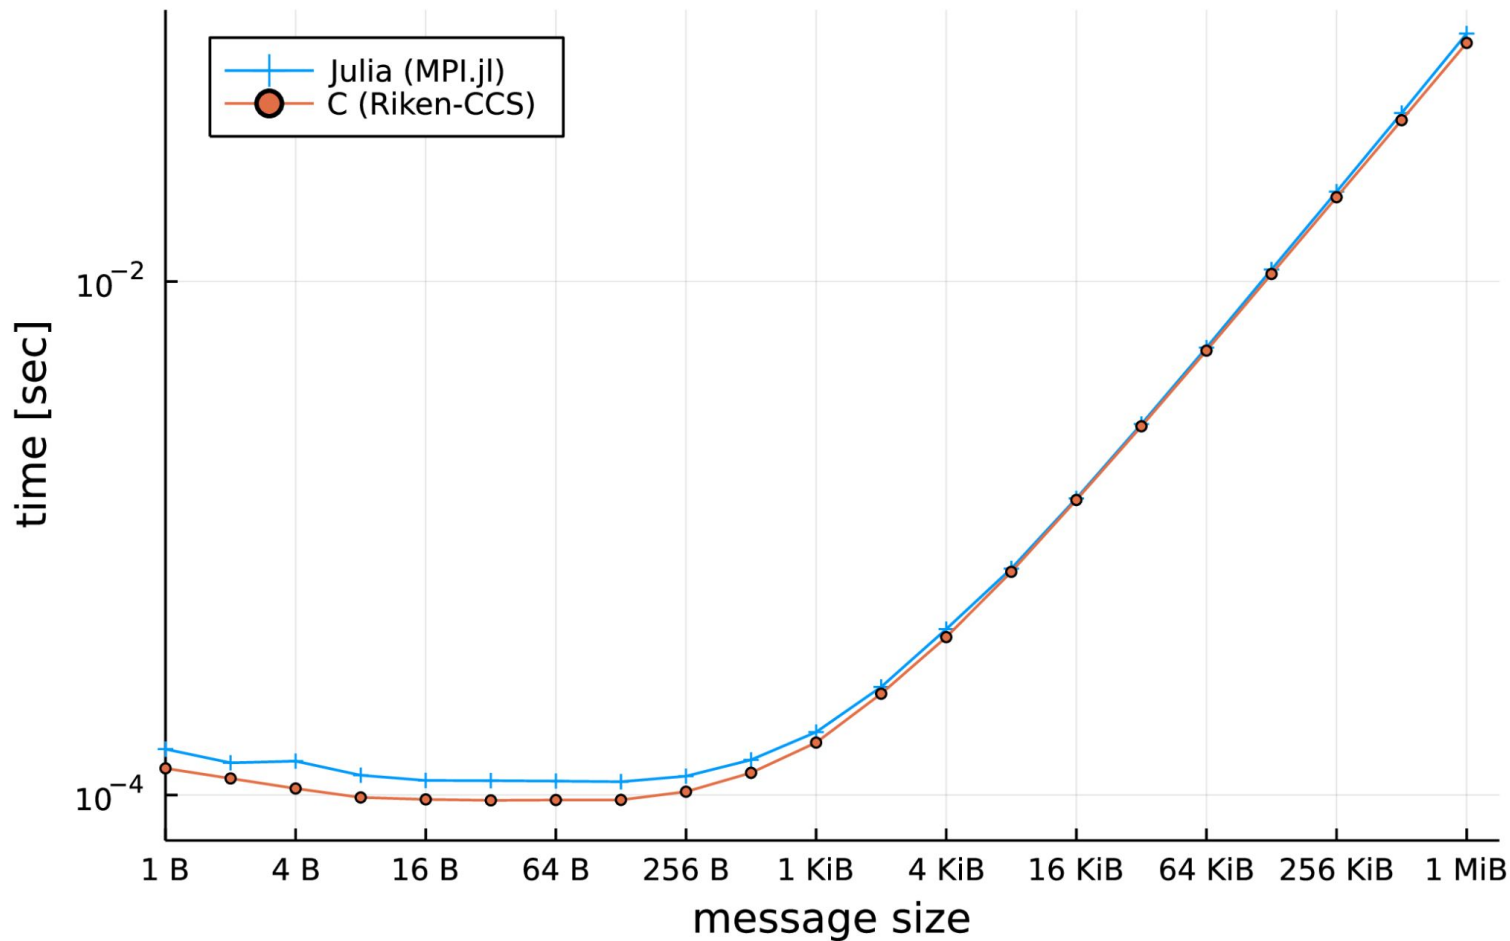
# Throughput of MPI PingPong @ Fugaku



Latency of MPI Allreduce @ Fugaku (384 nodes, 1536 ranks)



# Latency of MPI Gatherv @ Fugaku (384 nodes, 1536 ranks)



Latency of MPI Reduce @ Fugaku (384 nodes, 1536 ranks)

